

Certified Kubernetes Administrator Course

Introduction (S1)

- Info's about this course, the certificate and other courses
- Certificate:
 - 300\$ + free retry
 - Online examen
 - Mostly practice tasks

Core Concepts (S2)

Cluster Architecture

- https://github.com/BennyTheSen/kubernetes_absolute_beginners_course/blob/master/Kubernetes%20for%20beginners.pdf

ETCD for beginners

- Simple key value store
- Information stored as documents/pages for each entry (YAML, JSON, ...)
- `./etcdctl set <key> <value>`
- `./etcdctl get <key>`

ETCD in Kubernetes

- Stores information about Nodes, PODs, Configs, Secrets, ...
- Kubeadm does the installation for you and creates an etcd POD
- Multiple etcds should be connected to each other

Kube-API-Server

- Primary management component
- Kubectl commands work over the api server using POST and GET requests
- Responsible for:
 - Authenticate user
 - Validate request
 - Retrieve data
 - Update etcd
 - Scheduler
 - Kubelet
- Kubeadm deploys it as POD
 - Options under `/etc/kubernetes/manifests/kube-apiserver.yaml`
- *Non kubeadm setup:*
 - `/etc/sysdemd/system/kube-apiserver.service`
 - `ps -aux | grep kube-apiserver`

Kube Controller Manager

- Manages controllers in kubernetes
- Process that continuously monitors state of components of system and keeps system in desired state
- A lot of different types of controllers
 - All get installed with kubernetes Controller Manager
- Kubeadm deploys it as POD
 - Options under `/etc/kubernetes/manifests/kube-controller-manager.yaml`
- *Non kubeadm setup:*
 - `/etc/sysdemd/system/kube-controller-manager.service`
 - `ps -aux | grep kube-controller-manager`

Kube Scheduler

- Deciding which POD goes on which node and selects best Node
 - Depending on resource requirements, etc.
 - Can be customized
- Kubeadm deploys it as POD
 - Options under `/etc/kubernetes/manifests/kube-scheduler.yaml`
- *Non kubeadm setup:*
 - `/etc/sysdemd/system/kube-scheduler.service`
 - `ps -aux | grep kube-scheduler`

Kubelet

- create, delete containers on nodes with help of container runtime
- send information to master
- kubeadm does not install it automatically
 - `ps -aux | grep kubelet`

Kube Proxy

- service is virtual component only in kubernetes memory
- kubeproxy runs on each node and looks for new services and creates appropriate rules for network forwarding
- Kubeadm deploys it as POD on each node
 - Options under `/etc/kubernetes/manifests/kube-proxy.yaml`
- *Non kubeadm setup:*
 - `/etc/sysdemd/system/kube-proxy.service`
 - `ps -aux | grep kube-proxy`
-

POD's

- https://github.com/BennyTheSen/kubernetes_absolute_beginners_course/blob/master/Kubernetes%20for%20beginners.pdf

PODS with YAML

- https://github.com/BennyTheSen/kubernetes_absolute_beginners_course/blob/master/Kubernetes%20for%20beginners.pdf

Demo POD's with YAML

- https://github.com/BennyTheSen/kubernetes_absolute_beginners_course/tree/master/Section%206

Recap ReplicaSets

- https://github.com/BennyTheSen/kubernetes_absolute_beginners_course/blob/master/Kubernetes%20for%20beginners.pdf

Deployments

- https://github.com/BennyTheSen/kubernetes_absolute_beginners_course/blob/master/Kubernetes%20for%20beginners.pdf

Namespaces

- Every namespace has its own resources and policies
- To reach other namespace you have to append namespace to the pod name
- Kubectl get pods --namespace=<namespace>
 - Kubectl get pods --all-namespaces
- Kubectl create -f <yaml-file> --namespace=<namespace>
 - Or add namespace to metadata
- Create namespace with definition file and kubectl create
 - Or kubectl create namespace <namespace>
- Change default namespace:
 - Kubectl config set-context \$(kubectl config current-context) --namespace=<namespace>
- Resource quota to limit resources in namespace
 - With definition yaml file

Services

- https://github.com/BennyTheSen/kubernetes_absolute_beginners_course/blob/master/Kubernetes%20for%20beginners.pdf

Services Cluster IP

- https://github.com/BennyTheSen/kubernetes_absolute_beginners_course/blob/master/Kubernetes%20for%20beginners.pdf

Scheduling (S3)

Manual scheduling

- Scheduler searches for Pods without nodeName
 - Selects best node and sets the nodeName for the Pod
 - nodeName can be set manually in definition file
 - another way is with pod binding object (via definition yaml) and send POST request

Labels and Selectors

- labels: properties attached to items
 - added in definition file under metadata
- selectors: filters based on labels
 - *kubectl get <object-type> --selector <label>=<label-value>*
 - also used in replicaset, deployments, services, etc.
- annotations: additional information added to a object
 - added under metadata

Taints and Tolerations

- used to set restrictions to schedule which pods can go on which nodes
 - will not guarantee that a pod goes on a certain node
- Nodes have Taints
 - By default no taints on worker nodes
 - Master node has a default taint
 - *Kubectl taint nodes <node-name> key=<value>:<taint-effect>*
- Pods can have Tolerations
 - Default no tolerations
 - Specified in definition file under spec → tolerations

Node Selectors

- Limitation for pods to specify nodes
 - In pod definition file under spec > nodeSelector
 - Works with labels on nodes

Node Affinity

- To ensure pods are hosted on particular nodes
 - Specified under spec > affinity > nodeAffinity in pod definition
 - More complex than node selectors, also more options
- requiredDuringSchedulingIgnoredDuringExecution
 - node has to match or pod will not be scheduled
 - affinity ignored during execution
- preferredDuringSchedulingIgnoredDuringExecution
 - matching node is preferred but other nodes are used either
 - affinity ignored during execution
- requiredDuringSchedulingRequiredDuringExecution
 - node has to match or pod will not be scheduled
 - node has to match or pod will not be scheduled during execution

Taints and Tolerations vs Node Affinity

- for some cases taints and node affinity has to be used in combination

Resource Requirements and Limits

- scheduler takes resources in consideration for deploying pods
 - if not enough resources in any node POD does't get deployed
- Default Requests: 0.5 CPU (1=1vCPU), 256 Mi Memory, disk
 - Can be specified in deployment definition under spec > resources > requests
- Default Limits: 1 CPU, 512 Mi
 - Can be specified in deployment definition under spec > resources > limits
 - Limits can't be exceeded for CPU

DeamonSets

- Like replica sets but runs one copy of a POD on each node of the cluster
- Good for monitoring, logging, kube-proxy, networking, etc.
- Similar to replicaset definition just kind is different
- Uses default scheduler and nodeAffinity

Static Pods

- Nodes can be managed independently with kubelet
 - Pod definitions have to be placed under /etc/kubernetes/manifests
 - Path is option of the service pod-manifest-path or in kubeconfig.yaml
 - Only pods can be created like this
 - Kubelet tries to keep pods alive
 - Pods can be viewed with docker ps

Multiple Schedulers

- You can deploy your own scheduler
 - Can be specified with definition yaml too
- Cluster can have multiple schedulers
 - You can specify the scheduler when deploying a Pod
 - In the definition under spec > schedulerName
- View events with kubectl get events
- Kubectl logs <scheduler-name> --name-space<namespace>

Logging & Monitoring (S4)

Monitor Cluster Components

- Kubernetes doesn't have own monitoring solution
 - Other open source solutions like Prometheus, elastic stack, etc. work
 - Metrics Servers
 - As in memory solution
 - Kubelet cAdvisor retrieves performance metrics and exposes them to kubeApi
 - Get from github and deploy with kubectl create
 - *Kubectl top node* to view metrics of nodes
 - *Kubectl top pod*

Managing Application Logs

- Docker: docker logs [-f] <containerid>
- Kubernetes:
 - kubectl logs [-f] <pod-name>
 - with multiple containers: kubectl logs [-f] <pod-name> >containerid>

ALM (S5)

Rolling Updates and Rollbacks

- https://github.com/BennyTheSen/kubernetes_absolute_beginners_course/blob/master/Kubernetes%20for%20beginners.pdf

Commands

- Add commands with cmd or with a entrypoint or combination of both
 - FROM Ubuntu Entrypoint ["sleep"] CMD ["5"]

Commands and Arguments

- Add arguments under pod-definition containers>container>args
- Add commands under pod-definition containers>container>command

Configure Environment Variables in Application

- Set them under containers>container>env
 - List with name value pairs

Configure ConfigMaps in Application

- Used to pass config data from a file for env variables
 - Create config map and then inject in pod
 - *Kubectl create [-f] configmap [definition file]*
 - *Kubectl get configmaps*
 - *Kubectl describe configmaps*
 - Add to pod with envFrom

Configure Secrets in Application

- Used to store sensitive information
- Similar to config maps but get encoded
 - Encode data with `echo -n '<string to encode>' | base64`
 - Decode with `echo -n <encoded-string> --decode | base64`
 - *Kubectl create [-f] secret [definition file]*
 - *Kubectl get secrets*
 - *Kubectl get secret <secret> -o yaml*
 - *Kubectl describe secrets*
 - Add to POD with envFrom –secretRef or mount as Volume

Multi Container PODs

- https://github.com/BennyTheSen/kubernetes_absolute_beginners_course/blob/master/Kubernetes%20for%20beginners.pdf

Cluster Maintenance (S6)

OS Upgrades

- Pod eviction timeout... time until pod is considered dead if it not responds
 - *Kube-controller-mager --pod-eviction-timeout=5m0s*
- Drain nodes to get pods to other nodes *kubectl drain <node-name>*
 - *Kubectl uncordon <nodenname>* to remove drain
 - *Kubectl cordon <node-name>* to make node unschedulable

Kubernetes Software Versions

- See version with *kubectl get nodes*
 - Major.minor.patch, --alpha and --beta releases also available

Cluster Upgrade Process

- Scheduler and controller-manager can be a version lower than apiserver
- Kubelet and kube-proxy can be 2 versions lower than apiserver
- Kubectl can be a version higher or a version lower than apiserver
- Kubernetes only supports latest 3 versions
- Update one minor version at a time
- Update easy on cloud platform
- Kubeadm Update:
 - *Kubeadm upgrade plan*
 - *Apt-get upgrade --y kubeadm=<version>*
 - *kubeadm upgrade apply <version>*
 - *Apt-get upgrade --y kubelet=<version>*
 - *Systemctl restart kubelet*
- First upgrade master than the workers
 - Strategy1: Upgrade all at ones
 - Stretegy2: Upgrade one node at a time
 - *Kubectl drain <node>*
 - *Apt-get upgrade --y kubeadm=<version>*
 - *Apt-get upgrade --y kubelet=<version>*
 - *Kubeadm upgrade node config --kubelet-version <version>*
 - *Systemctl restart kubelet*
 - *Kubectl uncordon <node>*
 - Strategy3: add new nodes with new version, move workload and remove old node

Backup and Restore Methods

- Resource Configs:
 - Declarative configuration preferred way and store them in a src-repository
 - Backup all: *kubectl get all --all-namespaces -o yaml > all-deploy-services.yaml*
- ETCD-Cluster:
 - *Etcctl snapshot save <name>*
 - *Etcctl snapshot status <name>*
 - *Etcctl snapshot restore <name> --data-dir <directory>* afeter stopping kube-apiserver service, config service, reload daemon, restart etcd, restart kube-apiserver

Security (S7)

Kubernetes Security Primitives

- Hosts should be secured (no root access, pw based authentication disabled, only ssh)
- Kube-apiserver should be secured with authorization
 - Everything using tls encryption
- Restrict access between pods

Authentication

- Kubernetes does not handle users internally
 - Serviceaccounts can be managed with kubernetes
- Kube-apiserver auth mechanisms
 - Static password file
 - Eg. Csv. File with pw, username, userid, [group] and specify file in kubeapiserver service
 - Not recommended
 - Static token file
 - Instead pw use a token
 - Not recommended
 - Certificates
 - 3rd party identity services

TLS Basics

- Guarantee trust between parties in a transaction
- Asymmetric encryption(public key and private key) PKI(public key infrastructure)
 - Generate with ssh-keygen
 - On Server: openssl genrsa -out <key> <base>
 - Certificate to verify validity of key transfer (client and server certificates)
 - CA organizations sign certificates with their public and private keys (root certificates)
 - You can host private CA instances
 - Browsers check certificates
 - Server can request from client

TLS in Kubernetes

- Communication between nodes, components and users needs to be secured
- Servers:
 - Kube-apiserver exposes http service and requires server certificate
 - Etcd-server has requires own server certificate
 - Kubelet servers require server certificates too
- Clients:
 - Users like admin
 - Kube-scheduler
 - Kube-controller-manager
 - Kube-proxy
 - Kube-apiserver with etcd-server
 - Kube-apiserver with kubelet-server

TLS in Kubernetes – Certificate Creation

- Different tools like easyrsa, **openssl**, cfssl
- Kubeadm does most of the work for you
- CA Certificate:
 - Generate Key: `openssl genrsa -out <key-name> 2048`
 - Certificate Signing Request: `openssl req -new -key <keyname> -subj "/CN=KUBERNETES-CA" -out <keyfile>`
 - Sign Certificates: `openssl x509 -req -in <keyfile> -signkey <keyname> -out <certname>`
 - CA certificates files have to be well protected
- Client Certificates:
 - Admin user:
 - Like CA certificate with different names and "CN=kube-admin/O=system:masters"
 - `openssl x509 -req -in admin.csr -CA ca.crt -Cakey ca.key -out admin.crt`
 - Specify certificate in kube-config.yaml
 - Similar process for other client certificates
- Server Certificates
 - Etcd-server:
 - Similar to other certificates
 - Needs additional peer certificates
 - Kube-apiserver:
 - Create openssl.cnf with dns and IP's and path as option while generating
 - Kubelet-server:
 - Named after nodes and use them in kubelet-config.yaml
 - E.g.: Name: system:node:node01
 - Needs to be done for each node

View Certificate Details

- See spreadsheet in directory for what to check
- `Cat /etc/kubernetes/manifests/kube-apiserver.yaml`
- `openssl x509 -in <cert file> -text -noout` to view details
- inspect service logs for problems if kubeapi or etcd is down use docker logs

Certificates API

- Certificate signing request is send trough certificates API to CA server
 - Admin can approve request and share cert with user
 - Controller Manager does this
 - `openssl genrsa -out <key> 2048`
 - `openssl req -new -key <key> -subj "/CN=<username>" -out <file>`
 - Request specified with yaml file
 - `Kubectl get csr <request> -o yaml`
 - `Kubectl certificate approve <request>`

Kube Config

- Use kubeConfig file to store cert and key configurations and use the file for api requests
- 3 Sections:
 - Clusters: clusters you have access too
 - Users: users accounts you have you
 - Contexts: specify which user you want to use for which cluster
 - Namespace can also specified here
- *Kubectl config view [-- <config file>]*
- *Kubectl config use-context <user@cluster>*

API Groups

- Core group (/api):
 - Namespaces, pods, events, nodes, bindings, ...
- Named group (/apis):
 - /apps, /extensions, /networking.k8s.io, ... with their resources with their actions
 - See them under *curl <server-adress>:6443-k*
 - Use certificate files to authenticate
 - Or use *kubectl proxy*

Role Based Access Control

- Specified in role definition yaml
 - Role binding yaml object needs to be specified too
- *Kubectl get/describe roles*
- *Kubectl get/ describe rolebindings*
- Check access with:
 - *kubectl auth can-i <do something>*
 - *kubectl auth can-i <do something> as <user>*

Cluster Roles and Role Bindings

- see namespaced resources: *kubectl api-resources --namespaced=true*
- cluster roles/bindings for cluster scoped resources
- similar to non-cluster based roles
- can be used for namespace based resources too

Image Security

- hosting private image registry/registry is good for security
- *kubectl create secret docker-registry <name>*
 - specify secret under imagePullSecrets

Security Contexts

- container security settings can be specified in pod level
 - add securityContext under spec in definition file or under container

Network Policy

- PODs can communicate by default with each other
 - Can be prevented with network policies linked to PODs
 - Create with definition file and link with labels

Storage (S8)

Volumes

- In docker volume is linked to container to keep data persistent
- In kubernetes volume is attached to pod to store data
 - Specify volume in spec in definition file and use volume mounts
 - Also possible to specify different type of volumes like aws, gcp, ...

Persistent Volumes (PV)

- Cluster wide pool of storage volumes to manage storage more centrally
 - Users can select storage with volume claims
 - Create with a definition file
 - accessMode: Readonlymany, readwriteonce, readwritemany
 - Capacity
 - Volumetype: path, awsStore, etc.
 - *Kubectl get persistentvolume*

Persistent Volume Claims (PVC)

- Volumes getting bound to PVC
 - 1 to 1 relationship
 - If no volume available claim is pending state
- Specified with definition file
 - *Kubectl get persistentvolumeclaim*

Networking (S9)

Switching Routing

- Switching:
 - Switch connects host with interfaces in a network
 - *ip link*
- Routing:
 - Connect networks together
 - Router is also a device in the network
 - Gateway used as path to outside network
 - *route* to display routing table
 - *ip route add <ip> via <ip>*
 - *ip route add default via <ip>*
 - ip forwarding with */proc/sys/net/ipv4/ip_forward*
 - */etc/sysctl.conf* for persistent forwarding

DNS (Dynamic Name Service)

- Add known hosts to */etc/hosts* with *ip* and *hostname* (old way)
 - *Hostname* to find out the systems hostname
 - Works for commands like *ping*, *ssh*, *curl*, ...
- DNS Server resolves the names in modern day infrastructure
 - Add dns server to */etc/resolv.conf*
 - Add public nameservers like 8.8.8.8 for internet servers
 - *nslookup <dns-name>* to see server address or *dig*
- Domain Names:
 - Separated with dots to group websites under subdomains
 - *search* entry in */etc/resolv.conf* to look up for internal domain names

Network Namespaces

- processes in container have the same namespace
- host has own interfaces, routing tables, arp tables and can have its own network inside
 - *ip netns add* to add namespace
 - link two namespaces with:
 - *ip link add <veth name1> veth peer name <veth name2>*
 - *ip link set <namespace> netns <namespace1/2>*
 - *ip -n <namespace1/2> addr add <ip-addr1/2> dev <veth name1/2>*
 - *ip -n <namespace1/2> link set <veth name1/2> up*
 - to create virtual network between multiple networks, create a virtual switch
 - *ip link add <virtual network name> type bridge*
 - *ip link set dev <virtual network name> up*
 - *ip link add <veth name1> veth peer name <bridge network name1>*
 - *ip link set ip link set <veth name> netns <namespace>*
 - *ip link set <bridge network name1> master <virtual network name>*
 - *ip -n <namespace1/2> addr add <ip-addr1/2> dev <veth name1/2>*
 - *ip -n <namespace1/2> link set <veth name1/2> up*

Docker Networking

- container can be run with network option `--network`
 - none: no network connection to the outside
 - host: container is connected to host network
 - bridge: internal network is created, where containers are connected to

CNI (Container Network Interface)

- set of standards defining how programs should be developed to solve container networking
- already has a set of plugins available (BRIDGE, VLAN, IPVLAN, ...)
- also third party plugins weaveworks, flannel, etc.
- docker has its own standards

Cluster Networking

- kubernetes applications/parts listen to different ports (see documentation)

Pod Networking

- kubernetes doesn't have built in solution
 - every pod can reach every pod in same node with a unique ip address
 - every pod should be able to communicate with every pod on other nodes without NAT
- many solutions available like wavenet and flannel
- you can also build your own solution with some shell scripting and virtual networks

CNI in Kubernetes

- see *cni* with `ps -aux | grep kubelet` `ls /opt/cni/bin` and `/etc/cni/net.d/10-bridge.conf`

CNI weave works

- uses agents on every node for address resolution and ip address assignment
- can be deployed as service/deamon or as pod
 - `kubectl apply -f <weave_url with version>`
- `kubectl get pods -n kube-system`

IP Address Management – Weave

- Manual: host-local plugin for ip address management
 - can be configured in `net-script.conf`
- weave creates range from 10.32.0.1 to 10.47.255.254 on default

Service Networking

- clusterIP only reachable inside cluster
- exposes port to outside
- kube-proxy creates rules for iptables, ipvs, userspace
 - `cat /var/log/kube-proxy.log`

DNS in Kubernetes

- deploys built in dns server by default
- url for services: web-service.apps.svc.cluster.local

CoreDNS in Kubernetes

- kubernetes deploys dns server as kube-dns
 - core dns after 1.2.1
 - deployed as pod's in cluster which runs coredns executable
 - service is also deployed
 - */etc/coredns/Corefile*
 - *Kubectl get configmap*

Ingress

- Layer 7 Load-Balancer built-in kubernetes cluster
- Deploys solution(ingress controller) and specifies rules for configuration(ingress resources)
- Ingress controller:
 - Possible solutions: GCE, nginx, ...
 - Deployed as deployment with definition file, needs command to start
 - Pass in configmap object as easy configuration
 - POD_NAME and POD_Namespace as env variables
 - Also needs a NodePort service
 - Also a service account with roles, and rolebindings
- Ingress Resources:
 - Set of rules and configurations for the ingress controller
 - Forwarding, routing, ...
 - Created with yaml file as ingress object
 - Kubectl get/describe ingress
 - Configure rules in definition file under spec > rules > http > paths
 - Default backend for wrong routes
 - Use host field to specify the url

Install Hard Way (\$10)

Design a Kubernetes Cluster

- Purpose of Cluster?
 - Education: Minikube or Single Node cluster
 - Development & Testing: Multi Node Cluster with kubeadm
 - Production Application: Multi node cluster with multiple master nodes with kubeadm/GCP/AWS
- Cloud or on Prem?
 - Kubeadm for on prem
 - GKE for GCP, Kops for AWS, AKS for Azure
- Resources:
 - Storage:
 - high performance → SSD
 - multiple connections → network storage
 - also consider persistent storage
 - Nodes:
 - Virtual or physical machines
 - Master and worker nodes depending on workload
 - Not host workloads on master
- How many apps, what kind?
- Traffic amount

Choosing Kubernetes Infrastructure

- Not possible to install natively on windows → always running on Linux
- Minikube for single node cluster
- Kubeadm for single or multi-node cluster (requires VM's)
- Turnkey Solutions:
 - Provision, configure, maintain VM
 - Scripts for deployment
 - Kops on AWS
 - OpenShift: on prem solution with good ci-pipeline possibilities
 - CloudFoundry Container Runtime
 - VMware Cloud PKS
 - Vagrant
 - ...
- Hosted Solutions :
 - Kubernetes as a Service
 - Provider provisions, installs and maintains VM's
 - GKE on GCP
 - OpenShift Online
 - Azure Kubernetes Service
 - Amazon Elastic Container Service
 - ...

Choosing a Network Solution

- <https://kubernetes.io/docs/concepts/cluster-administration/networking/#how-to-implement-the-kubernetes-networking-model>

Configure High Availability

- Consider use of multiple master nodes
- Also redundancy of all other control plane components
 - Api-server: can be active on all master nodes, node balancer should be used
 - Scheduler and Controller manager: should run in active-standby mode
 - Needs leader elect option in kube-controller manager
 - ETCD: can be on master node or external(safer version)

ETCD in HA

- Distributed database which ensures consistency
 - One node is leader with RAFT protocol
 - Write operation is complete if it is written in quorum(>50%) of nodes
 - Odd number of nodes should be preferred
- Download etcd binary → install → configure service
- etcdctl for commands
 - export ETCDCTL_API=3
 - etcdctl put <key> <value>
 - etcdctl get <key>
 - etcdctl get / --prefix --keys-only (for all keys)
- etcd nodes for HA setup is 3,6 or 7

"Kubernetes the Hard Way" on Virtualbox

- <https://github.com/mmumshad/kubernetes-the-hard-way>

TLS Bootstrap Worker Nodes

- Certificate operations done automatically with kube-api
- Permissions on master node have to be made
- Needs a bootstrap token for permissions
- System:node-bootstrapper is default role
- Kubectl get csr to view certificate requests
- Can be automatically approved by role certificatesigningrequests:nodeclient
- Automatic approvement only for client certificates possible

Install Kubeadm Way (S11)

- https://github.com/BennyTheSen/kubernetes_absolute_beginners_course/blob/master/Kubernetes%20for%20beginners.pdf
- <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>

End to End Tests on a Kubernetes Cluster (S14)

End to End Tests

- Manual:
 - Kubectl get nodes
 - Kubectl get pods -n all-namespaces
 - Service Status
 - Kubectl run nginx
 - Kubectl scale
 - Kubectl expose ...
- Test suite: <https://github.com/kubernetes/test-infra>
 - Around 1000 tests spread across 7 different areas (takes ~ 12 hours)
 - Creates namespace → creates object → executes curl on object → records result → cleanup
 - Also good to build and test own kubernetes based solution (needs to pass conformance tests)
- Sonobouy is another tool

End to End Tests – Run and Analyze

- go get -u k8s.io/test-infra/kubetest
- kubetest -extract=<kubernetes-version>
- cd kubernetes
- kubetest -test -provider=<provider> <outfile>
 - skeleton as local provider
 - - test_args to focus on a specific area

Smoke tests

- Demo for cluster installed in S10

Ent to End test part 1

- Demo for cluster installed in S10

Troubleshooting (S13)

Application Failure

- Check accessibility
 - Map your application services and connections and check all the parts
 - App: curl
 - Service: kubectl get/describe/logs service
 - Pod: kubectl get/describe/logs pod [--previous]
- <https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/>

Control Plane Failure

- Kubectl get nodes
- Kubectl get pods --all-namespaces
- Service kube-apiserver/kube-controller-manager/kube-scheduler/kubelet/kube-proxy status
- Kubectl logs kube-apiserver

Worker Node Failure

- Kubectl get nodes
- Kubectl describe node
- Top
- Df -h
- Service kubelet status
- Sudo journalctl -u kubelet
- Check certificates

Network Failure

- No content atm

Other Topics (\$14)

JSON PATH

- <https://kodekloud.com/p/json-path-quiz>

Advanced Kubectl Commands

- Kubectl commands work with JSON
 - Kubectl commands can be extended with json path
 - `kubectl get pods -o json`
 - `kubectl get pods -o=jsonpath='{ <query> }'`
 - `kubectl get pods -o=custom-columns=<Colum name>:<jsonpath>`
 - `kubect get nodes --sort-by=<json path query>`